

## Setting up VirtualBox Environment

1. Create following Networks in VirtualBox:
  - # Host-Only Ethernet Adapter #2, Configure Manually, IPv4 Address: 10.0.0.1, IPv4 Network Mask: 255.255.255.0, no DHCP
  - # NAT Network "ProviderNetwork1", Network CIDR: 203.0.113.0/24
2. Create following nodes in VirtualBox:
  - # Name: "controller"
    - i. 2 VCPU
    - ii. 6144 MB RAM
    - iii. 20GB Primary Disk
    - iv. Network Interfaces:
      - Adapter 1: connected to NAT
      - Adapter 2: connected to Host-Only Network #2
      - Adapter 3: connected to NAT Network "ProviderNetwork", enable Promiscuous Mode
  - # Name: "compute1"
    - i. 1 VCPU
    - ii. 2048 MB RAM
    - iii. 10GB Primary Disk
    - iv. Network Interfaces:
      - Adapter 1: connected to NAT
      - Adapter 2: connected to Host-Only Network #2
      - Adapter 3: connected to NAT Network "ProviderNetwork", enable Promiscuous Mode
  - # Name: "compute2"
    - i. 1 VCPU
    - ii. 2048 MB RAM
    - iii. 10GB Primary Disk
    - iv. Network Interfaces:
      - Adapter 1: connected to NAT
      - Adapter 2: connected to Host-Only Network #2
      - Adapter 3: connected to NAT Network "ProviderNetwork", enable Promiscuous Mode
  - # Name: "block1"
    - i. 1 VCPU
    - ii. 1024 MB RAM
    - iii. 10GB Primary Disk + Second Disk 30GB
    - iv. Network Interfaces:
      - Adapter 1: connected to NAT
      - Adapter 2: connected to Host-Only Network #2
      - Adapter 3: connected to NAT Network "ProviderNetwork", enable Promiscuous Mode
  - # Name: "deployment"
    - i. 1 VCPU
    - ii. 2048 MB RAM

- iii. 40GB Disk
- iv. Network Interfaces:
  - Adapter 1: connected to NAT
  - Adapter 2: connected to Host-Only Network #2

## Ubuntu Linux Installation on Target VMs

3. Install Ubuntu 16.04 Server in Minimal VM configuration on each node.

## Preparing Target VMs

4. On Each Node, configure GRUB to enable standard Network Interface Names:
  - # Edit /etc/default/grub and modify the GRUB\_CMDLINE\_LINUX to "net.ifnames=0 biosdevname=0"
  - # Run "update-grub" as superuser
  - # Run "visudo" to enable primary user to "sudo su" without password:  
<username> ALL=(ALL) NOPASSWD:ALL
  - # Set up /etc/hosts:  
127.0.0.1 localhost  
10.0.0.11 controller  
10.0.0.31 compute1  
10.0.0.32 compute2  
10.0.0.41 block1  
10.0.0.100 deployment
5. On **controller**, **compute1** and **compute 2** set up /etc/network/interfaces:  

```
auto lo
iface lo inet loopback

auto eth0
iface eth0 inet dhcp

auto eth1
iface eth1 inet static
    address <node ip address>
    netmask 255.255.255.0
    dns-nameservers 8.8.8.8

auto eth2
iface eth2 inet manual
    up ip link set dev eth2 up
    down ip link set dev eth2 down
```
6. On **controller** run following commands as superuser:
  - # apt update -y
  - # apt upgrade -y
  - # apt install -y python python-simplejson glances vim
  - # reboot

7. On **compute1** and **compute2** run following commands as superuser:

```
# apt update -y
# apt upgrade -y
# apt install python python-simplejson glances vim
# echo "configfs" >> /etc/modules
# update-initramfs -u
# reboot
```

8. On **block1** set up `/etc/network/interfaces`:

```
auto lo
iface lo inet loopback

auto eth0
iface eth0 inet dhcp

auto eth1
iface eth1 inet static
address 10.0.0.41
netmask 255.255.255.0
dns-nameservers 8.8.8.8
```

9. On **block1** run following commands as superuser:

```
# apt update -y
# apt upgrade -y
# apt install -y python python-simplejson glances vim
# apt install -y lvm2 thin-provisioning-tools
# pvcreate /dev/sdb
# vgcreate cinder-volumes /dev/sdb
# echo "configfs" >> /etc/modules
# update-initramfs -u
# reboot
```

## Prepare deployment VM

1. Configure GRUB to enable standard Network Interface Names:

```
# Edit /etc/default/grub and modify the GRUB_CMDLINE_LINUX to "net.ifnames=0 biosdevname=0"
# Run "update-grub" as superuser
# Run "visudo" to enable primary user to "sudo su" without password:
<username> ALL=(ALL) NOPASSWD:ALL
# Set up /etc/hosts:
127.0.0.1 localhost
10.0.0.11 controller
10.0.0.31 compute1
10.0.0.32 compute2
10.0.0.41 block1
10.0.0.100 deployment
```

**2. Set up /etc/network/interfaces:**

```
auto lo
iface lo inet loopback

auto eth0
iface eth0 inet dhcp

auto eth1
iface eth1 inet static
    address 10.0.0.100
    netmask 255.255.255.0
    dns-nameservers 8.8.8.8
```

**3. Run following commands as superuser:**

```
# apt update -y
# apt install -y python-jinja2 python-pip libssl-dev curl
# pip install -U pip
# apt upgrade -y
# reboot
```

**4. Configure & test ssh access to other nodes:**

```
# ssh-keygen -t rsa
# ssh-copy-id <username>@controller
# ssh-copy-id <username>@compute1
# ssh-copy-id <username>@compute2
# ssh-copy-id <username>@block1
# ssh <username>@controller
# ssh <username>@compute1
# ssh <username>@compute2
# ssh <username>@block1
```

**5. Verify network configuration:**

```
# ifconfig
# ip a
```

**6. Install Ansible and Kolla-Ansible:**

```
# pip install ansible==2.5.2
# pip install kolla-ansible==6.0.0
# cp -r /usr/local/share/kolla-ansible/etc_examples/kolla
/etc/kolla
```

**7. Prepare nova-compute.conf file with QEMU configuration:**

```
# mkdir -p /etc/kolla/config/nova
# vim /etc/kolla/config/nova/nova-compute.conf
[libvirt]
virt_type = qemu
cpu_mode = none
```

## Kolla-Ansible Parameters

1. Edit `/etc/kolla/globals.yml` file:  
config\_strategy: "COPY\_ALWAYS"  
kolla\_base\_distro: "centos"  
kolla\_install\_type: "binary"  
kolla\_release: "pike"  
kolla\_internal\_vip\_address: "10.0.0.10"  
network\_interface: "eth1"  
neutron\_external\_interface: "eth2"  
neutron\_plugin\_agent: "openvswitch"  
keepalived\_virtual\_router\_id: "51"  
nova\_console: "novnc"  
enable\_cinder: "yes"  
enable\_cinder\_backup: "no"  
enable\_cinder\_backend\_iscsi: "yes"  
enable\_cinder\_backend\_lvm: "yes"  
enable\_ha\_proxy: "yes"  
enable\_heat: "yes"  
enable\_horizon: "yes"  
enable\_openvswitch: "{{ neutron\_plugin\_agent != 'linuxbridge' }}"  
keystone\_token\_provider: "fernet"  
fernet\_token\_expiry: 86400  
glance\_backend\_file: "yes"  
cinder\_volume\_group: "cinder-volumes"  
nova\_compute\_virt\_type: "qemu"
2. Edit multimode file:  

```
# cp /usr/local/share/kolla-ansible/ansible/inventory/* .
# vim multimode
[control]
controller ansible_ssh_user=<username>
ansible_become=True
ansible_private_key_file=/root/.ssh/id_rsa
[network]
controller ansible_ssh_user=<username>
ansible_become=True
ansible_private_key_file=/root/.ssh/id_rsa
[compute]
computel          ansible_ssh_user=<username>
ansible_become=True
ansible_private_key_file=/root/.ssh/id_rsa
compute2          ansible_ssh_user=<username>
ansible_become=True
ansible_private_key_file=/root/.ssh/id_rsa
[monitoring]
```

```
controller ansible_ssh_user=<username>
ansible_become=True
ansible_private_key_file=/root/.ssh/id_rsa
[storage]
blockl          ansible_ssh_user=<username>
ansible_become=True
ansible_private_key_file=/root/.ssh/id_rsa
```

## Running Kolla-Ansible Deployment

1. Generate Kolla-Ansible passwords:  
# kolla-genpwd
2. Bootstrap Nodes:  
# kolla-ansible -i multinode bootstrap-servers
3. Run pre-deployment checks:  
# kolla-ansible -i multinode prechecks
4. Run Deployment:  
# kolla-ansible -i multinode deploy
5. Verify that containers are up, not restarting:  
# docker ps -a

## Finalize Deployment

1. Run Kolla Post Deploy script:  
# kolla-ansible post-deploy
2. Install OpenStack Client:  
# pip install python-openstackclient
3. Populate OpenStack with initial set up:  
# vim /usr/local/share/kolla-ansible/init-runonce  
# . /etc/kolla/admin-openrc.sh  
# cd /usr/local/share/kolla-ansible  
# ./init-runonce

## Verification and Using the OpenStack

1. Check the **admin** password:  
# grep keystone\_admin\_password /etc/kolla/passwords.yml
2. Open browser and go to 10.0.0.10, login as admin